

All Step One-Pagers

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

July 2026

All Step One-Pagers

Contents

Step 1 One-Pager — Reinforcement Learning Basics.....	3
Step 2 One-Pager — Game Theory and CFR Basics.....	4
Step 3 One-Pager — CFR Variants and Monte Carlo Methods.....	5
Step 4 One-Pager — Game Abstraction and Scaling.....	6
Step 5 One-Pager — Neural Networks for Imperfect-Information Games.....	7
Step 6 One-Pager — End-to-End Game AI Architectures.....	8
Step 7 One-Pager — Opponent Modeling.....	9
Step 8 One-Pager — Safe Exploitation.....	10
Step 9 One-Pager — Multi-Agent Reinforcement Learning.....	11
Step 10 One-Pager — Population-Based Training and Evolutionary Game Theory.....	12
Step 11 One-Pager — Dynamic Coalition Formation in Competitive FFA Games.....	13
Step 12 One-Pager — Sequence Models and LLM Agents in Strategic Settings.....	14

Step 1 One-Pager — Reinforcement Learning Basics

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

April 2026

Problem. Everything later in this plan — neural equilibrium finding (Step 5), the opponent-model actuator (Steps 7-8), multi-agent learning (Step 9) — is assembled out of value networks, policy gradients and experience replay. Step 1 is the foundation layer: how an agent improves a policy from reward alone, in the fully observable single-agent setting, before that picture is complicated by a second, hidden, adversarial player. It feeds Contribution #1 indirectly (the architecture patterns are reused by the adaptive opponent model) and directly supplies the implementation vocabulary every later step assumes.

Approach. One algorithm from each family, written from scratch in PyTorch: **DQN** (value-based, off-policy) on **CartPole-v1** — Q-network [4->128->128->2], circular replay buffer, frozen target network synced every 5 episodes — and **PPO** (policy-gradient, on-policy) on **LunarLander-v3** — separate actor [8->128->128->4] and critic, GAE ($\lambda=0.95$) by reverse sweep, clipped surrogate at $\epsilon=0.2$. Each was then benchmarked against its **Stable-Baselines3** counterpart at matched core hyperparameters and matched step budgets (DQN 750K, PPO 500K), scored on best rolling-100 average so that early stopping cannot flatter the result. **All numbers below are measured.**

Key results (measured).

- *Both targets met.* DQN solves CartPole at **episode 1011**, 100-episode average **477.5** (target 475); PPO solves LunarLander at **264,192 steps** at **202.2** (target 200) — 236K steps inside its budget.
- *The from-scratch versions beat the SB3 defaults on these tasks.* Best rolling-100: DQN **477.5 vs 293.9**, PPO **203.6 vs 131.2**. The cause is algorithmic, not unfair tuning: episode-based epsilon decay against SB3's step-based one, an adaptive target sync (5 episodes is ~100 steps early and ~2500 once solved) against a fixed 1000 steps, and early stopping against running the full budget.
- *eps_min was the decisive DQN knob.* At 0.01 the run peaked at **479.1** (episode 810) and decayed to **302.6**; 0.001 plus best-model checkpointing turned a degrading policy into a stable one.
- *Capacity, not budget, was PPO's bottleneck.* [64, 64] peaked at **179.9** and never crossed 200; [128, 128] cleared it. Advantage normalisation proved non-negotiable on rewards spanning -500 to +300, and without the 0.01 entropy bonus the policy collapsed to hovering in place.
- *Honest read on the comparison.* SB3's defaults are tuned for robustness across environments, not for classic control; its PPO was still climbing at 500K and would likely reach target given 1-2M steps. The claim is a task-specific advantage, not a better algorithm.

Thesis connection. The 300-500-line stack (against SB3's ~10K) is the deliberate price paid for surgical modifiability: injecting belief-state observations into an actor-critic is a deep subclassing exercise in SB3 and a local edit here, which is exactly what the exploitation work in Steps 7-8 needs. The same actor-critic pair returns as MAPPO/MADDPG in Step 9, and the value-network machinery returns as Deep CFR's advantage network in Step 5.

Open questions. The MDP assumption that makes all of this work — a fully observable state — is precisely what poker breaks: over an information set, DQN's **max** and PPO's per-state policy are both ill-posed, which is why Step 2 changes tool rather than scale. Open: how much of the single-agent stability toolkit (target networks, trust regions, advantage normalisation) survives when the environment is itself a learning agent (deferred to Step 9's non-stationarity), and whether the sample-efficiency edge measured here is a real property or an artifact of two small, well-conditioned control tasks.

Step 2 One-Pager — Game Theory and CFR Basics

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

April 2026

Problem. Step 1’s machinery assumes a fully observable state and a stationary environment; poker has neither. Two-player zero-sum imperfect-information games need a different solution concept — Nash equilibrium defined over *information sets* — and a different quality measure: exploitability, not reward. Step 2 establishes both, plus the algorithm that computes them, because everything after it is either an approximation of CFR (Steps 3-5), a deliberate departure from the equilibrium CFR produces (Steps 7-8), or an attempt to say what equilibrium even means once $N > 2$ (Steps 9, 11).

Approach. Vanilla CFR written from scratch on **Kuhn Poker** — 3 cards, 2 players, 12 information sets: small enough to have a closed-form equilibrium *family* (parameterised by α in $[0, 1/3]$), rich enough to contain bluffing, mixed strategies and indifference. Hand-coded components: the game engine, recursive counterfactual traversal with regret matching and chance sampling, and an exact exploitability evaluator that enumerates all $2^6 = 64$ pure strategies — because a per-state “oracle” best response is *not* a best response; the best-responder must commit to one action per information set. Trained for 100,000 iterations, with an OpenSpiel cross-verification script alongside. **All numbers below are measured** (implementation/step02/models/cfr_results.json).

Key results (measured).

- *CFR lands inside the analytical equilibrium family, not merely near it.* The recovered bluff parameter is $\alpha = 0.1941$, and the family’s own internal constraint $P(\text{bet} \mid K) = 3 \cdot \alpha$ holds to **2.6e-4** (measured **0.58247** against the **0.58221** its own α implies).
- *Pure decisions converge hard; mixed frequencies carry the $O(1/\sqrt{T})$ tail.* King bets and calls at ≥ 0.9999 , Jack folds to a bet at **0.99998**, Queen passes at the root at **0.99993** — while the mixing sits **0.002-0.011** off the closed form (Jack bluff after a pass **0.3403** vs $1/3$; Queen call **0.5387** vs $1/3 + \alpha = 0.5274$). The frequencies are the slow part, which is precisely where the residual regret lives.
- *Game value -0.0602* measured against the exact $-1/18 = -0.0556$ at 100k iterations, reproducing Player 0’s structural first-mover disadvantage.
- *The convergence rate is the theoretical one.* Log-log exploitability slope **-0.489** against a predicted **-0.5**.
- *Game value alone is not a validity check.* A strategy that always bluffs the Jack can still average near $-1/18$ while being trivially exploitable; only exploitability catches it. That is why exploitability — not reward — is the metric carried forward into Steps 7-8 and 14.

Thesis connection. Nash is the baseline this thesis exists to leave: Contribution #1 reads a specific opponent in order to justify departing from it, and Contribution #2 bounds how far the departure may go. Concretely, the exact best-response and exploitability code written here becomes the literal oracle reused in Steps 7-10, and Kuhn remains the smallest testbed where every later claim can be bracketed by an exact answer.

Open questions. Why the *average* strategy converges when the current strategy does not — the intuition (averaging damps the overshoot, as Polyak averaging does) is not a proof. Everything here rests on the two-player zero-sum minimax theorem, and neither CFR’s guarantee nor exploitability’s meaning survives into N-player or general-sum settings (Steps 9, 11). And full-tree traversal touches every information set on every iteration, so this exact algorithm is already out of budget one game up — the motivation for Step 3’s Monte Carlo sampling and Step 4’s abstraction.

Step 3 One-Pager — CFR Variants and Monte Carlo Methods

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

April 2026

Problem. Step 2’s vanilla CFR visits every information set on every iteration, which stops being affordable one game up from Kuhn. The field’s two answers are a better-constant full-traversal method (**CFR+**) and *sampling* the tree (**MCCFR**), and the received wisdom is that sampling is what made real poker solvable. Step 3 asks which one actually wins, on what size of game, and why — the question that decides the solver every later step depends on.

Approach. All four algorithms hand-coded in Python + NumPy on **Leduc Poker** (6 cards, two betting rounds, a revealed community card, 936 information sets, 10,200 nodes, 120 deals): vanilla CFR, CFR+ (regret flooring, linear averaging, alternating updates), MCCFR **external sampling** (all traverser actions, one sampled opponent action), and MCCFR **outcome sampling** (one root-to-terminal trajectory, **eps** = 0.6 on-policy mixture with importance-sampling correction). Graded by an exact information-set-constrained best response, under a common **180-second wall-clock budget**, with OpenSpiel as the reference solver. **All numbers below are measured.**

Key results (measured).

- *CFR+ is the workhorse, and it is not close.* At essentially the same iteration count as vanilla CFR (**3,706 vs 3,713**), CFR+ reaches **2.6e-5** exploitability against vanilla’s **4.4e-3** — roughly **170x** better, from three localised code changes.
- *MCCFR loses badly at this game size.* External sampling reaches **5.5e-2** after **3.5M** iterations and outcome sampling **1.0e-1** after **8.3M** — three to four orders of magnitude worse than full traversal despite ~1,000x more iterations.
- *The reason is variance, not speed, and it is quantified.* Sampling is **945x / 2,228x** cheaper per iteration but needs **~99,000x / ~520,000x** more of them. Netted into a wall-clock constant $C_w = C/\text{sqrt}(\text{speed})$: **0.075** (vanilla) against **0.767** and **1.143** — sampling is **10.2x / 15.3x** slower to any given accuracy.
- *The crossover is derivable and Leduc is far below it.* Sampling breaks even at roughly **2.1M nodes** (external) and **4.8M** (outcome); Leduc’s 10,200 nodes are **210x / 466x** too small, while limit Hold’em (~1e14) is orders of magnitude past it. This reconciles “MCCFR was essential for real poker” with “MCCFR loses here” — same algorithm, opposite side of the line.
- *Cross-validated against OpenSpiel* at a matched 500-iteration budget: vanilla **0.020 vs 0.022**, CFR+ **8.6e-4 vs 9.4e-4**.
- *One measured anomaly, unexplained.* Vanilla CFR beats its own worst-case rate: the supposedly constant $\text{eps} \cdot \text{sqrt}(T)$ falls from **0.90** ($T=100$) to **0.27** ($T=3,700$) instead of holding flat, so the $O(1/\text{sqrt}(T))$ bound is loose here in a way the bound itself does not predict.

Thesis connection. This step fixes the tooling for everything that follows: Leduc plus the exact best-response evaluator become the standard testbed of Steps 7-12, and CFR+ becomes the Nash reference against which every later opponent-modelling, safe-exploitation and LLM number is bracketed. The exactness-versus-variance trade-off measured here returns in Step 5, where the sampled estimates feed a network instead of a table.

Open questions. Whether the crossover node count is a property of the algorithms or of this implementation’s constant factors — a vectorised or compiled full-traversal solver shifts **speed_v** and moves the threshold, so the number is honest for this code and not yet a claim about the algorithms in general. Which sampling scheme survives once the strategy is a neural network rather than a table (Step 5). And why vanilla CFR converges faster than its worst case.

Step 4 One-Pager — Game Abstraction and Scaling

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

May 2026

Problem. Step 3 established that a solver’s reach is bounded by tree size, so the classical route to real poker is not a better solver but a *smaller game*: shrink it, solve the shrunken version, translate the strategy back, and pay an exploitability tax for the distinctions you threw away. Step 4 measures that tax. The question is not “does abstraction help” but which axis of compression is safe, and whether the compute a reduction buys ever repays the strategic information it destroys.

Approach. A Leduc-family pipeline that compresses along three axes and evaluates every resulting strategy **in its own full game**: **lossless suit isomorphism** (Leduc payoffs depend only on rank and on whether the private card pairs the board, so suit-isomorphic information sets merge with no loss), **lossy card bucketing** (hand-strength features, EMD-style distances, k-means, in perfect- and imperfect-recall regimes), and **action abstraction plus translation** (nearest-action, probability-split, pseudo-harmonic) on a variable-bet Mini-NL Leduc. Benchmarked with CFR+ under a common **180-second** budget, 3 seeds, on fixed-limit Leduc, Mini-NL Leduc and a 4-rank Extended Leduc. **All numbers below are measured.**

Key results (measured).

- *Lossless compression is the highest-value move, and it wins by buying iterations.* Suit isomorphism cuts fixed-limit Leduc from **936 to 288** information sets and reaches **3.13e-6** exploitability against full CFR+’s **4.44e-5** in the same wall clock — because it completes **14,185** iterations instead of **2,655**. Confirmed at scale on Extended Leduc: **10,304 -> 2,968** info sets (-71%), **0.0272 -> 0.00126**, roughly 6x the iterations.
- *Lossy card buckets install a floor that compute cannot lift.* **k2 0.571, k3 0.382, k5 0.382** — all at 11,000+ iterations, i.e. 4x more solving than the full game and four orders of magnitude worse. The error is in the abstraction, not the budget.
- *Action abstraction is the dangerous axis.* Mini-NL: the full 4,704-info-set game reaches **0.00692**, while the 936-info-set action abstraction reaches **0.673** despite 5x the iterations — roughly **100x worse** for a 5x smaller tree. Extended Leduc is starker still: suit+action **4.696** and suit+action+buckets **4.734**, both far worse than simply solving the unabridged game. Translation error dominates the total.
- *A measured artifact, not a law.* **k5** matches **k3** exactly because Leduc’s post-flop hand-strength distributions collapse into three effective shapes; the bucket-count ordering seen here is a property of this tiny game.
- *Stated limitations.* CFR+ is deterministic, so the three seeds are not independent stochastic runs; the action-abstraction numbers depend on the current translator and deployment semantics and should be read as a diagnostic failure mode; and the OpenSpiel comparison aligns all 936 information sets as a sanity check, not as proof of identical solver dynamics.

Thesis connection. The reporting format is the transferable part: strategy quality is meaningless without game size beside it, and the Pareto frontier of the two is an early seed of Contribution #3’s evaluation methodology. The mechanism transfers as well — Step 7’s type-based opponent modelling is this same “group for tractability” idea moved from state space to strategy space, and its confident-but-wrong failure is action translation’s misspecification in another costume.

Open questions. Static translation is brittle by construction, which points at subgame and nested solving (Step 6) as the production-grade answer to off-tree actions. Whether the ordering measured here — card abstraction survivable, action abstraction ruinous — holds in a game with enough hand-strength diversity that **k5** genuinely differs from **k3**. And whether an abstraction can be adapted to the opponent rather than fixed in advance, which is where this step touches the adaptive framework directly.

Step 5 One-Pager — Neural Networks for Imperfect-Information Games

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

May 2026

Problem. Steps 2-4 all store a strategy as a table indexed by information set, and Step 4 showed both escapes from that table — a better solver, a smaller game — running out. The third escape is to stop enumerating: replace the regret table with a function approximator that *generalises* across similar information states, at fixed memory. Step 5 studies that substitution — Deep CFR, its single-network variants, NFSP — and what it costs in guarantees.

Approach. Scope, stated honestly: **this step stopped after exploration and reading — there is no from-scratch Phase-4 implementation and no implementation report**, and the survey’s closing synthesis is still a stub. The measured work is a two-day exploration against OpenSpiel’s reference solvers on Leduc and Kuhn — Deep CFR at three network sizes, NFSP on both games, tabular MCCFR and CFR+ as baselines — graded by OpenSpiel’s own exploitability so the head-to-head is honest (Step 3’s MCCFR runs on a custom engine and is not comparable). **All numbers below are measured** (implementation/step05/exploration/logs/day0{1,2}_results.json).

Key results (measured).

- *The step’s most actionable artifact is a library bug.* OpenSpiel 1.6.12’s PyTorch DeepCFRSolver **never trains its advantage networks**: if `len(samples.info_state == 0)` takes `len` of an elementwise comparison, always truthy, so it returns before the optimiser step. Advantage losses come back silently `None` and exploitability sits at random-strategy level (**~1.69**) whatever the iteration count. A one-character fix, patched locally.
- *Where the table fits, the table wins — decisively.* Leduc at comparable wall time: CFR+ **0.00123** (400 iter, 92 s), tabular MCCFR **0.0967** (50k iter, 68 s), Deep CFR **1.49-1.70** (120 iter, ~95 s), NFSP **2.46** (50k episodes) — CFR+ roughly **1,400x** better than Deep CFR.
- *Those neural numbers are budget artifacts, not verdicts.* Deep CFR’s curves are flat from iteration 30 to 120 where the paper uses 400+, and OpenSpiel’s own NFSP Leduc example runs **2e7** episodes — **400x** our budget. Rule adopted: distrust an NFSP Leduc curve below ~1e6 episodes.
- *Smaller network won at this budget.* (32,32) **1.1441** beat (64,64) **1.7002** and (128,128,128) **1.4908** — consistent with underfitting on few distinct samples, and expected to reverse with more iterations.
- *What under-training looks like inside the policy.* Against a 400-iteration CFR+ reference, Deep CFR at 40 iterations has **median total-variation 0.35** across 8 sampled Leduc info states (range 0.163-0.708): the direction of each decision is usually right, but the probabilities are pulled toward uniform (**0.73/0.27** where CFR+ is **0.92/0.07**) — MSE regression on noisy counterfactual samples buys an under-confident smoothing of equilibrium.
- *The cost structure that justifies the family anyway.* A Deep CFR outer iteration costs **~0.7 s** against a tabular MCCFR iteration’s **~1.4 ms** (~500x), but its per-iteration work is roughly constant in game size while the tabular cost scales with information-set count. Leduc is a teaching benchmark, where neural methods lose to baselines they were never meant to beat.

Thesis connection. Deep CFR’s advantage network is Step 1’s value network with a counterfactual target — the joint that lets Steps 6-8 leave table-sized games behind. More pointedly, the under-confident smoothing measured here is the failure mode Step 7 later prices: a model with the right direction and the wrong frequency is what makes best-responding to an imperfect read lose to Nash.

Open questions. The step’s own gate is unmet: a hand-coded Deep CFR with a Vitter reservoir sampler, and DREAM’s outcome sampling with baseline subtraction, both remain outstanding, so “can I build it” is unanswered. Untested: whether the tabular-versus-neural ordering flips exactly at Step 3’s node-count crossover, and whether the info-state tensor’s structure really matters more than network depth.

Step 6 One-Pager — End-to-End Game AI Architectures

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

June 2026

Problem. Steps 2-5 assembled parts in isolation — CFR variants, abstraction, neural value approximation — without asking how they compose into a system that beats humans, or what such a system guarantees. Step 6 surveys the five competition-grade architectures of 2017-2023 as an evolution of design trade-offs rather than a leaderboard: the state of the art this dissertation builds on and departs from.

Approach. A deliberately *architectural* study — what each system computes offline, what it computes in real time, where learning sits, what it can prove — of **DeepStack** (2017), **Libratus** (2017/18), **Pluribus** (2019), **ReBeL** (2020) and **Student of Games** (2023), each scored on the same nine dimensions and placed on three axes: abstraction to neural representation, offline precomputation to real-time search, imperfect-information-only to unified play. **This step produced no code; every number below is the result each paper reports, not a measurement of ours**, and the chapter itself is drafted but not yet signed off.

Key results (as reported by the papers).

- *The abstraction era was far more exploitable than its match results suggested.* A local-best-response probe showed top competition bots losing **over 3,000 mbb/g** — several times the win rate that made them look strong. Closing that gap is what the decade was for.
- *Search at inference, not the offline strategy, carries the edge.* Libratus’s raw blueprint **lost** to Baby Tartanian8 by **8 mbb/g**; with nested safe subgame solving on, the same system **won by 63**, and beat four professionals by **147 mbb/g over 120,000 hands**. Solving off-menu bets live also beat nearest-size action translation **119 vs 1,465 mbb/g** worst-case — Step 4’s translation failure repaired by refusing to translate.
- *Cost is not the axis it looks like.* DeepStack spent **~175 CPU-core-years** labelling values and Libratus **~25M core-hours**, while Pluribus’s blueprint cost **~12,400 core-hours** — about **\$150** on one server — and still beat five elite professionals at **+48 mbb/g**.
- *Every advance was paid for elsewhere.* Pluribus reached six players by discarding **all** safety guarantees for unsafe search; ReBeL recovered the two-player guarantee and eliminated abstraction and blueprint alike (Dong Kim **+165 mbb/g**) but retreated to two players; Student of Games unified perfect- and imperfect-information play soundly (beating the LBR probe by **+434 mbb/g**) while sitting **over 1,100 Elo** below Go specialists — its authors’ “price of generality”. No system dominates every axis.
- *The one omission all five share.* Each is **opponent-blind by design**: it computes a worst-case-robust strategy and plays it unconditionally. The human-tested systems say so explicitly, and Pluribus does not even know who it is playing.

Thesis connection. That shared omission is the opening this dissertation occupies, and the chapter supplies the machinery to fill it. ReBeL’s **public belief state** is the substrate Contribution #1 widens from beliefs over cards to beliefs over opponent *type*, inheriting the soundness proofs attached to it rather than bolting adaptation on outside. Pluribus is Contribution #2’s argument in one system: if equilibrium buys no safety at $N > 2$ anyway, declining to exploit forfeits a guarantee never held. Exploitability, the LBR probe (repurposed from certificate to stress test) and AIVAT variance reduction are three ready instruments for Contribution #3.

Open questions. Four, straight from the synthesis: opponent-blindness itself (Step 7); safe exploitation beyond two-player zero-sum, which Pluribus won without and Student of Games could not extend (Step 8); carrying an opponent model *past the depth limit* instead of discarding it at the leaf; and real-time compute budgets, whose six-order-of-magnitude spread binds hardest on an agent that must re-solve *and* re-estimate an opponent at once. Corollary: build on the cheap, open lineage — the flagship code is unreleased or supercomputer-scale.

Step 7 One-Pager — Opponent Modeling

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

July 2026

Problem. A Nash-equilibrium strategy is unexploitable but blind: it plays identically against everyone and never punishes a weak opponent. Opponent modeling is the sensor that turns observed actions into an estimate of a specific opponent’s strategy, so the agent can deviate from equilibrium to exploit them — the first half of the thesis’s Behavioral Adaptation Framework (Contribution #1). The hard part: doing so without becoming exploitable yourself.

Approach. Bayesian belief-update loop (prior x likelihood -> posterior -> best response), built behind one shared interface so a single exact best response can be applied to three interchangeable models: (1) **type-based** — a belief over a fixed menu of opponent types; (2) **continuous** — free-form per-situation Dirichlet counts; (3) **consistent** — a sequence-form convex-optimization estimate (Ganzfried 2025). An **adaptive exploiter** runs the observe -> model -> best-respond -> act loop, with a Nash safety blend and an optional change-point detector for non-stationary opponents. Testbeds: Kuhn Poker and Leduc Hold’em, both exactly solvable, so every result is bracketed by exact analytical references.

Key results (measured).

- *Modeling is worth it.* Equilibrium play leaves **0.11-0.28 per hand** on the table against exploitable Kuhn opponents; the gap against a Nash opponent is ~ 0 (you cannot exploit an equilibrium).
- *When the model class fits, best response reaches the exact ceiling* — the type-based model is statistically indistinguishable from the best-response ceiling on every opponent in both games.
- *A confident-but-underfit model makes you exploitable.* On Leduc the continuous model **loses to Nash (-0.175 vs a -0.083 ceiling)** — best-responding to a wrong estimate of an unexploitable opponent opens a leak in its own play. This is the exploitation-vs-safety tension in a single number.
- *Confident-and-wrong is a real risk.* Against an out-of-menu opponent, the type-based posterior commits hard to the nearest representable type; over 300 seeds a wrong type led past hand 100 in $\sim 13\%$ of runs (it always corrects eventually, and never falls once locked).
- *Non-stationarity is scenario-dependent.* Change-point forgetting rescues a harmful stale model (Kuhn: $-0.116 \rightarrow +0.226$) but underperforms simple adaptation when the new opponent is exploitable and the detector false-fires (Leduc). The reaction matters as much as the detection.
- *Consistency.* The consistent model recovers Kuhn strategies accurately (TV $\sim 0.004-0.021$) but its per-update convex solve is too costly for the online loop as built — an empirical answer to the step’s real-time-feasibility question, pointing to incremental solving / Step 8 approximations.

Thesis connection. Step 7 builds the **sensor** (the model); Step 8 builds the **actuator** (safe, KL-regularized exploitation). The continuous model’s Nash self-leak is the empirical case for that safety mechanism; the consistency theory is the principled backbone the framework extends.

Open questions. Real-time consistent modeling (incremental convex solve); principled non-stationarity (confidence-scaled forgetting, better change signals); out-of-menu opponents (open-world type discovery); one opponent to many (joint best response is not the combination of individual ones, and the convex guarantee breaks for $N > 2$).

Step 8 One-Pager — Safe Exploitation

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

July 2026

Problem. Step 7 built the *sensor* — an opponent model — and showed its danger: best-responding to an imperfect model can open a leak in your own play (the continuous model *lost* to Nash on Leduc). Step 8 builds the *actuator*: turn a model into profit **without becoming exploitable**. This is the second half of the thesis’s Behavioral Adaptation Framework (Contribution #1) and the launch point for multi-agent safe exploitation (Contribution #2).

Approach. Every safe-exploitation method is the same program — *maximize expected value against the opponent model, subject to a safety floor on your worst-case value* — which is linear in the hero’s sequence-form realization plan. One LP engine solves it by **constraint generation** (solve $\rightarrow$ read policy $\rightarrow$ call an exact best response as the worst-case oracle $\rightarrow$ add a cut if unsafe $\rightarrow$ re-solve), reusing Step 7’s engines, best-response, Nash and opponent zoo. Five families differ only in the floor: **RNR** (tunable p), **Ganzfried** ($\geq$ Nash value v^*), **prime-safe** ($\geq v^* - \epsilon$ for an ϵ -equilibrium baseline), **SES** (blueprint value, enforced locally on a subgame via a gadget), **adaptation** (no more exploitable than the blueprint). Testbeds: Kuhn and Leduc, both exactly solvable, so every result is bracketed by exact analytical references. **All numbers below are measured.**

Key results (measured).

- *Safe exploitation works on Kuhn.* **Ganzfried is safe against every opponent** (worst-case $\geq v^*$ within $1e-3$) **and beats Nash’s own EV on every exploitable type** — e.g. **+0.222 vs Nash’s +0.146** against **AlwaysPass**. Full best response earns more (+0.975) but its worst-case collapses to **−0.5** (ruinously exploitable). This is the central validated result.
- ϵ *is measured, not fabricated.* Prime-safe/adaptation lower the floor to $v^* - \epsilon$ with the baseline’s exploitability measured at $\epsilon = 0.0074$, and earn a little more than Ganzfried by spending that budget (+0.266 vs +0.222 on **AlwaysPass**).
- *Canonical RNR is bang-bang, not a smooth frontier* (a contradicted prediction). In a game this small the max-min LP jumps from the safe vertex straight to full best response at $p \approx 0.7$ — the smooth curve is the *naive blend*’s, and it is dominated at the safe corner.
- *Headline — global safe-exploitation does not scale; local does.* On **Leduc**, the global solvers (Ganzfried/prime-safe/adaptation) **fail to converge within a 40-iteration cap** (worst-case **−0.64 to −1.33**, grossly unsafe), while the **subgame method (SES) converges** (194–350 iters) and stays near-safe (worst-case $\approx$ **−0.13**) while extracting **+0.25 to +0.68** vs weak types. This is the global-vs-local safety theory $\rightarrow$ practice gap, measured on a tiny game.
- *The safety guarantee bites against a worst-case adversary, not a benign one.* In the teaching attack (bait $\rightarrow$ Nash reveal), the honest separating signal is the safety-violation count (**full_br 40/40 refits, Ganzfried 0/40**), not realized profit — a gentle Nash “revealer” never claws back full_br’s bait-phase windfall. A punishing test needs an adaptive counter-exploiter.

Thesis connection. Step 7 = sensor; Step 8 = actuator. The Kuhn results confirm the actuator is sound; the Leduc non-convergence is the concrete argument for real-time subgame methods (SES / OX-Search) and/or an exact dual-LP formulation, and it is the empirical bridge into the scalable, multi-agent safe exploitation of Contribution #2.

Open questions. Scalable safety (exact dual-LP vs local/subgame — the Leduc wall); the SES gadget’s residual exploitability ($0.04 > \text{tol}$ — provably or only approximately safe?); a punishing teaching attack (adaptive reveal); and **N-player safety** — every guarantee here rests on the two-player zero-sum fact that Nash secures v^* against any opponent, an anchor that vanishes for $N > 2$ (Contribution #2).

Step 9 One-Pager — Multi-Agent Reinforcement Learning

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

July 2026

Problem. Steps 2–8 lived inside two-player zero-sum games, where a Nash strategy secures a value v^* against anyone and CFR provably converges to it. Step 9 is the pivot into the multi-agent world, where the defining difficulty is **non-stationarity**: every agent is learning at once, so from any one agent’s seat the environment (which contains the others) never holds still. This is the launch point for the thesis’s multi-agent contributions — dynamic opponent modeling (#1), safe exploitation without a minimax anchor (#2), and population-level evaluation (#3).

Approach. Build and compare the field’s four structural answers to non-stationarity on small, *exactly-solvable* testbeds, so every learner is graded against ground truth. **Independent learning** (the control that fails) on four canonical matrix games; **CTDE** — centralize the critic at training, decentralize the actor at execution — via MADDPG and MAPPO on cooperative tasks; **PSRO** — a game played over a *population* of policies (meta-Nash + best-response oracle), reusing Step 07’s exact best response as both oracle and exploitability metric — on Kuhn, Leduc, a matrix game, and a native Goofspiel; **learned communication** (CommNet); and **LOLA** (differentiate through the opponent’s learning step) on the Iterated Prisoner’s Dilemma. **All numbers below are measured.**

Key results (measured).

- *Independent learning fails exactly where theory says.* It converges to Nash on Prisoner’s Dilemma (defection), Stag Hunt (the risk-dominant corner), and Battle of the Sexes, but **Matching Pennies never converges** — its distance-to-Nash actually *grows* ($0.30 \rightarrow 0.48$), a sharper-than-predicted non-convergence (a contradicted prediction, §9 of the report).
- *PSRO is the game-theory $\leftrightarrow$ MARL bridge, and it works on the small games.* Exploitability $\rightarrow$ machine zero on **Kuhn** ($0.917 \rightarrow 2 \times 10^{-16}$, 6 rounds), $\rightarrow 0$ on the matrix game and Rock–Paper–Scissors. Self-play confirms the mechanism: on Kuhn the **average**-iterate NashConv falls $0.24 \rightarrow 0.031$ while the **last** iterate keeps oscillating — why a population/averaging is needed.
- *A centralized critic is a near-zero-variance teacher.* On CoopSignal the centralized critic’s residual is $3.2e-11$ vs the independent critic’s 0.077 — the CTDE variance-reduction claim, confirmed.
- *Communication clears the guessing ceiling.* CommNet with the channel ON scores 0.795 vs 0.204 OFF (ceiling $1/K = 0.2$) — a learned, not designed, protocol.
- *LOLA turns defectors into cooperators.* IPD per-step return rises from 1.04 (naive defection) to 2.82 (LOLA cooperation); zeroing the look-ahead recovers the naive gradient exactly (the mechanism check).
- *Honest negatives (kept predictions, §9).* PSRO on **Leduc** declines but hits a scaling wall ($4.75 \rightarrow 2.16$ after 20 rounds, not < 0.5); **Goofspiel K=4** oscillates rather than converging (a flagged code anomaly, documented not fixed); and on the **climbing game** no method reaches the optimum 11 (IL/MAPPO 7, MADDPG 5) — a centralized critic lowers variance but does not by itself solve hard-exploration coordination.

Thesis connection. PSRO is Step 2’s iterated best response lifted to a population and the empirical backbone here; LOLA is *dynamic* opponent modeling, the moving-target complement to Step 7’s static read (Contribution #1); PSRO’s meta-game is a population-level evaluation methodology (Contribution #3). The Leduc scaling wall echoes Step 8’s global-vs-local finding, and the place where every two-player guarantee stops — the **missing N>2 minimax anchor** — is exactly Contribution #2’s problem statement.

Open questions. Can an approximate (RL) oracle push Leduc PSRO close to Nash, and at what cost to the guarantee? What breaks the Goofspiel-K=4 oscillation (de-duplication? a mixed population? a different meta-solver?), and why does a centralized critic *hurt* on the climbing game? And underneath all of it: with no minimax value for $N>2$, what does “safe” mean for a population, and can PSRO’s meta-game supply the substitute (Contribution #2)?

Step 10 One-Pager — Population-Based Training and Evolutionary Game Theory

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

July 2026

Problem. Self-play can go in circles: in a cyclic game (rock-paper-scissors) “get better” has no meaning, so a population that only trains against itself spins instead of improving. Step 10 is the population-level layer of the thesis — how to *diagnose* whether a game/population is a skill ladder or a wheel of counters, and how to *train and evaluate* a whole population of agents. It is the launch point for automated opponent modeling (#1), population safe-exploitation without a minimax anchor (#2), and population-level evaluation (#3).

Approach. Two halves, both graded against exact references. **Part I — evolutionary game theory as a diagnostic:** replicator dynamics on four solvable matrix games (checked against analytic ESS/Nash) and the transitive/cyclic **spinning-top** decomposition. **Part II — an AlphaStar-style PBT league** of neural PPO agents on Leduc Hold'em (three agent types: main / main-exploiter / league-exploiter, plus freezing and PFSP), evaluated with **EGTA/meta-Nash**, Elo, and diversity metrics — every neural policy extracted to a *tabular* policy so Step 07's **exact** best response measures its exploitability. **All numbers below are measured.**

Key results (measured).

- *Replicator dynamics reproduce every analytic outcome.* Prisoner's Dilemma → all-Defect; Hawk-Dove → the interior 0.5 ESS (orbit radius 0.0); **Rock-Paper-Scissors never converges** (orbits the centre, radius 0.095); Stag Hunt → a basin-dependent pure ESS.
- *The spinning-top diagnostic works, and Leduc's structure depends on the population.* Hodge decomposition gives RPS transitive 0.0 / cyclic 1.0 and a skill ladder 1.0 / 0.0 (the SVD rank-1 method wrongly gives RPS 0.707). The **PSRO best-response** meta-game on Leduc is mostly **cyclic** (≈ 0.45 transitive, 27 three-cycles); the **league snapshot** meta-game is mostly **transitive** (≈ 0.94 -0.98) — a contradicted prediction, §4/§7 of the report.
- *The league produces strong individuals.* Its best frozen snapshot reaches exploitability 1.305 (scale), beating exact PSRO (2.163) and self-play (3.683); CFR-Nash floor is 0.0099.
- *Honest negatives (kept predictions).* League exploitability is **non-monotone at scale** ($4.73 \rightarrow \min \approx 1.21 \rightarrow 2.05$ over 120 epochs) — the live agents regress late; the best agents are frozen snapshots. And the **meta-Nash mixture is more exploitable than its best member** at scale ($3.418 > 1.305$) — meta-Nash minimizes meta-game regret, not full-game exploitability, so mixing behavioral policies can hurt. Diversity is thin (participation ratio 1.9, a single behavioral cluster).

Thesis connection. The league is Step 9's PSRO made asynchronous with neural oracles, reusing Step 07's exact best response as the yardstick; main exploiters are automated opponent modelers (Contribution #1); EGTA/meta-Nash is the population evaluation methodology (Contribution #3). The league's missing guarantee — it can regress, and its mixture can be exploitable — is the population form of the **missing $N > 2$ safety anchor** (Contribution #2). The transitive/cyclic diagnostic predicts Step 11's FFA coalition games will be strongly cyclic.

Open questions. Does best-snapshot retention / population regularization remove the late regression (a training fix) or is it inherent to the three-type design (a design fix)? Should a population ship a selected best-response-robust member or a diversity-regularized meta-solver instead of the meta-Nash mixture? And underneath both: with no minimax value for a population, what does “safe” mean, and can the exploiter mechanism be given a guarantee rather than staying heuristic (Contribution #2)?

Step 11 One-Pager — Dynamic Coalition Formation in Competitive FFA Games

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

July 2026

Problem. Steps 2-10 all leaned on one crutch: a two-player game with an exact best-response and exploitability oracle, so “did it work?” was one number. Add a third player and alliances become possible — form, exploit, betray — while Nash turns both intractable and strategically empty, since it ignores coalitions entirely. Step 11 removes the crutch deliberately. This is the thesis frontier, not a consolidation step.

Approach. A native 4-player **So Long Sucker** engine — the 1950 game Nash, Shapley, Shubik and Hausner built to study alliances — implemented from the De Carufel & Jerade formalization, with a 2-player **minimax endgame** as its only exact anchor. On top: a **coalition detector** (help/harm from chip placement), **Shapley credit** redefined as the value of a member’s win probability, a **coalition-aware MAPPO** trainer blending coalition and winner-takes-all reward by a weight **alpha**, and **EGTA + spinning-top** analysis reused from Step 10. **All numbers below are measured**, with one caveat: `scale_results.json` is a **pre-fix** run cited only as evidence of the bug below — the authoritative figures are `sweep_scale.json` (5 seeds) and post-fix `smoke_results.json`.

Key results (measured).

- *Certified wherever exactness exists.* 100/100 random games terminate, all rewards are zero-sum, `endgame_mismatches`: 0 against exact minimax. The detector recovers a **planted {0,1} alliance** from the move stream alone (pair score **10.0**, cross-pair entries 0 or -1), and Shapley reproduces the glove and majority toys exactly, core emptiness included.
- *The step’s biggest finding is a bug, and its lesson outlived it.* Two red checks shared one cause: **~99.5% of random games end in deadlock**, so the winner falls to `_most_chips`, whose lowest-index tie-break gave seat 0 twice its fair share (all-random winners [94,42,33,31]). With an unbiased tie-break, symmetric Shapley spread fell **0.525** -> **0.013** and winners went uniform; the same fix deflated an impressive **0.87** hero win rate to an honest **~0.41** against a 0.25 floor. In a game that nearly always ends near-tied, the tie-break rule is load-bearing.
- *“Coalitions don’t emerge at scale” was overturned by finding the right knob.* **alpha** dominates and the **0.3 default sits in a dead zone**: at **alpha ~ 0** the proxy-credit gap is **+0.0376 +/- 0.0103** (~4.4x the sparse **0.0109**), while **every alpha >= 0.3 cell is negative**. The effect *grows* with game size (~10x smoke), and the cheap critic-value proxy beats the expensive counterfactual credit (**+0.038 vs +0.013**).
- *Coalition behaviour is bought, not free.* **alpha = 0** drops hero win rate to **~0.29**, essentially the random floor, while **alpha >= 0.1** holds **~0.52**.
- *Strongly cyclic, honestly short of the bar.* The skill-ladder pool is transitive-dominant (cyclic **0.25-0.31**), but a coalition pool pushes cyclic to **~0.57-0.69** — large, yet still failing the strict “>50% dominance” rule, so the check stays red. Harness: **4/5 PASS**.

Thesis connection. Contribution #1 generalises cleanly: opponent modelling lifts from “what kind of player is this” to “who is allied with whom”, read straight off the moves. Contribution #2 gets its problem statement sharpened rather than solved — with an **empty core** and no Nash anchor, “safe” cannot mean bounded deviation from equilibrium and must become behavioural/population-based (piKL). Contribution #3 gains a working substitute for exploitability in EGTA plus the spinning-top decomposition, inheriting Step 10’s caveat that the population you assemble decides whether you see a ladder or a wheel.

Open questions. Reconciling the engine’s turn and tie-break model against De Carufel & Jerade is the highest-value follow-up — both the seat-0 bug and the red cyclic check trace to documented simplifications, and the theorems remain an unverified reading item. Whether a 3- or 4-player-aware EGTA projection surfaces the cycling the 2-type collapse discards. And what “safe” can mean at all when the core is empty.

Step 12 One-Pager — Sequence Models and LLM Agents in Strategic Settings

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

July 2026

Problem. Two post-classical ways to play a game bypass everything Steps 2-11 built. **Sequence modelling** recasts play as conditional prediction — feed a transformer (**return-to-go**, **state**, **action**) triples and ask for the next action given a desired return, no value function, no self-play — and **ARDT** adds minimax relabelling for worst-case robustness. **LLM agents** skip training entirely: hand a model the rules in English. Step 12 scores both against an exact benchmark.

Approach. Kuhn Poker, where Step 2 supplies exact Nash and exact exploitability, with a narrow port to Leduc Hold'em as a complexity check. Compared: Nash-CFR, behavioural cloning, a Decision Transformer, ARDT, and four LLM backends (offline stub, gpt-oss-20b, Qwen2.5-7B, OpenThinker3-7B), plus follow-ons asking *why* that ranking falls out. **All numbers below are measured** from committed JSON, with two caveats: two results were **retracted** mid-session (an impossible +1.59 gap closed, and a Leduc -0.83 from a decoder discarding 70% of the probability mass), and the Nash reference is 0.0162 chips at 5,000 CFR iterations but 0.0061 at 50,000 — both ~0.

Key results (measured).

- *The step's two nominal subjects finish last.* **BC 0.055 « ARDT 0.469 < DT 0.799 < LLM 0.833**, against Nash **0.0162**. Cloning lands within 0.04 chips of Nash by copying a near-Nash policy, while the DT routes that same policy through a return-conditioning channel carrying mostly luck — conditioning actively destroys information here.
- *Return conditioning never steers, on either game.* **high(+2) = 0.6981 vs low(-2) = 0.6928**: tied, and ordered wrong. The Kuhn explanation (a 4-value payoff alphabet) was **half-refuted** on Leduc: 15 payoff values made the predicted artefact vanish, but steering still failed to appear ($r = +0.062$).
- *The models play better than they can explain.* The hypothesis going in was “knows the right frequency, cannot sample it”; the data says the reverse. Scored as strategies, *stated* frequencies are far worse than *played* ones — Qwen **0.921 vs 0.357** chips, gpt-oss **1.576 vs 0.392**. Probe behaviour; do not ask.
- *Exploitative by default, not adaptive.* Mean gap closed **+0.38** but mean **learning -0.22**: against the one well-powered opponent the model captures **83%** of available exploitation from the first half onward and never improves. It exploits **61% harder than Nash while being 59x more exploitable**, and only against passive opponents.
- *Scale is irrelevant, and the LLM edge belongs to Kuhn.* **Qwen-7B beats gpt-oss-20B by +0.162 chips/hand** over 20k hands per pair, strictly transitively. On Leduc the LLM is indistinguishable from the DT (-0.463 vs -0.454) and cannot beat weak opponents at all (-0.071 where Nash gets +0.582), with **100%** of its illegal-action mass a single misconception: folding when checking is free.
- *One scalar hides the diagnosis.* A single Queen node is **41.4%** of total exploitability, while value-betting the King at 1.00 where Nash mixes at 0.68 costs **0.1%** — deviation size and cost are nearly uncorrelated. Harness: **3 PASS / 2 FAIL**, honestly red.

Thesis connection. Contribution #1 gets a negative result worth having: behavioural adaptation is *absent* in-context, so an explicit opponent model must be built rather than assumed. Contribution #2 gets its frontier measured on one plot — 61% more exploitation for 59x the exploitability, only against weak opposition. Contribution #3 gets the per-decision decomposition and the measurement protocol.

Open questions. The named fix is ARDT's $Q(s,a)$: a state-only return estimator cannot tell “this state is bad” from “*this action* is bad”, which is why the tau sweep contradicts theory and exploitability comes out *lowest* on the optimistic side — the default was left at the theoretically correct value rather than switched to buy a better number. Whether the Leduc misconception is one prompt line away. And the standing caveat: every Leduc LLM conclusion rests on one model, one prompt style, 600 hands.