

Step 7 Summary — Opponent Modeling in Imperfect-Information Games

Research on the possibilities for applying Artificial Intelligence in computer games

Alexander Andreev

July 2026

Contents

7	Step 7 — Opponent Modeling in Imperfect-Information Games	1
7.1	1. Why Opponent Modeling — the money Nash leaves on the table	2
7.2	2. The Bayesian Core — belief, evidence, response	3
7.3	3. Seeing Through a Keyhole — the partial-observability problem	5
7.4	4. Three Models, One Interface	7
7.5	5. The Consistency Problem and the Sequence-Form Fix	8
7.6	6. When a Model Is Confident and Wrong	10
7.7	7. From Model to Money — best response, the ceiling, and a self-inflicted leak	12
7.8	8. Adapting to Change — non-stationary opponents	14
7.9	9. Connections and Forward Pointers	15

7 Step 7 — Opponent Modeling in Imperfect-Information Games

This is a ground-up chapter on opponent modeling: the problem, the mathematics, the family of methods, and a set of controlled experiments run on two small poker games. It serves two purposes — a self-contained refresher while later steps build on it, and a primary source for the Step 15 public synthesis. It is written to be read on its own; no prior familiarity with the project’s code is assumed. All experimental numbers reported here were measured on reproducible runs of the two testbeds (Kuhn Poker and Leduc Hold’em) and are bounded, wherever possible, by *exact* analytical references rather than simulated ones.

Where this sits in the thesis. A game-theoretic agent from the earlier steps computes a *Nash equilibrium* — a strategy that cannot be beaten in the long run. This step builds the opposite capability: a **sensor** that watches how a *specific* opponent actually plays and turns those observations into an estimate of their strategy, so the agent can deviate from equilibrium to punish that opponent’s mistakes. That sensor is the first half of the thesis’s Behavioral Adaptation Framework (Contribution #1). The second half — how to exploit *safely*, without opening a hole in your own play — is formalized in Step 8.

7.1 1. Why Opponent Modeling — the money Nash leaves on the table

A **Nash equilibrium** strategy is built never to lose in the long run, *no matter who it plays*. In a two-player zero-sum game that guarantee is real and valuable: an equilibrium strategy has a fixed *value*, and no opponent — however clever — can push you below it. But the same property is also a limitation. An equilibrium strategy plays **identically** against a world champion and against someone who folds every single time you bet. It never adapts, because adaptation is exactly what it was designed not to need.

That refusal to adapt has a price, and the price is measurable. **Opponent modeling** is the act of watching a specific opponent’s actions, forming a belief about their strategy, and **deviating from Nash to exploit the pattern** — bluffing more against someone who folds too much, value-betting thinner against someone who calls too much.

The cleanest way to see the value is a *rock-paper-scissors* picture. Your opponent secretly throws rock 70% of the time. The safe strategy is to randomize evenly and break even forever. But if you *notice* the bias, you throw paper more and start winning. Two things make this the whole step in miniature: you must **infer** the bias from a noisy stream of throws (a few aren’t enough), and if you **over-commit** to paper you become predictable and they crush you with scissors. Inferring the bias is opponent modeling; knowing how far to lean is the exploitation-versus-safety tradeoff that runs through everything below.

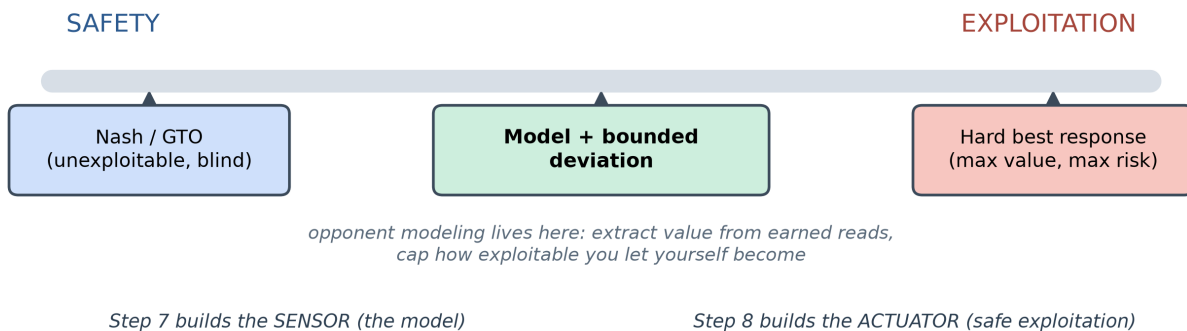


Figure 1: The safety-exploitation dial: pure equilibrium play is unexploitable but blind; a hard best response extracts the most value but is maximally risky. Opponent modeling operates in between. Step 7 builds the sensor (the model); Step 8 builds the actuator (safe exploitation).

7.1.1 How much is at stake — measured on Kuhn Poker

Kuhn Poker is the smallest non-trivial poker game: a three-card deck (Jack, Queen, King), one card each, a single betting round. It is small enough to solve *exactly*, which lets us compute the value of modeling rather than guess at it. For each of several fixed opponent styles we computed three exact quantities:

- **Nash EV** — what an equilibrium strategy earns against that opponent (per hand);
- **Best-response EV** — what a strategy that *knows the opponent perfectly* earns: the exact ceiling on exploitation;

- **the gap** between them — the money equilibrium play leaves on the table.

Opponent (Kuhn)	Nash EV	Best-response EV	Exploitation gap
Calling station (never folds)	+0.119	+0.333	0.215
Rock (only commits with the King)	-0.060	+0.167	0.227
Maniac (always bets/calls)	+0.121	+0.333	0.213
Nash (equilibrium play)	-0.055	-0.052	0.004

Against the three exploitable styles the gap is enormous — **0.11 to 0.28 per hand**, on a game whose entire equilibrium value is about one twentieth of a chip. Against a Nash opponent the gap is essentially **zero**, exactly as theory demands: you cannot exploit an equilibrium, because best response cannot beat the game value. (The small negative Nash EVs are Kuhn’s known first-player disadvantage of $-1/18 \approx -0.056$ — that is the game, not an error.)

Two lessons are already visible. First, **modeling is worth doing**: the value is large and real. Second, **exploitation is directional** — the exact best response to a rock *bluffs more* (it bets the weakest hand into an open pot, because the rock folds everything but the King). The rest of this chapter is about earning that value without giving it back.

Read more: Southey, F. et al. (2005). “Bayes’ Bluff: Opponent Modelling in Poker.” *UAI*.

7.2 2. The Bayesian Core — belief, evidence, response

Underneath every method in this chapter is one loop, and it is worth stating in plain language before any notation. You are a detective. You start with a hunch about how the opponent plays (a **prior**). Each action they take is a clue: it makes some explanations more likely and others less. You fold the clue into your suspicion (the **posterior**), and repeat. Opponent modeling *is* this loop:

belief about their strategy $\rightarrow$ see an action $\rightarrow$ update belief $\rightarrow$ best-respond $\rightarrow$ repeat.

Formally, the update is Bayes’ rule. If σ denotes a candidate opponent strategy and a the action just observed,

$$P(\sigma \mid a) \propto P(\sigma) \cdot P(a \mid \sigma),$$

in words: **posterior** $\propto$ **prior** $\times$ **how well that strategy explains what they just did**. An action a candidate calls impossible drives its probability toward zero; an action it predicts well boosts it. Over many hands the belief concentrates on the best-fitting explanation.




Figure 2: The opponent-modeling loop: a prior belief is multiplied by the likelihood of each observed action to give a posterior, which drives a best response; the outcome feeds the next observation.

7.2.1 The type “zoo”

The simplest way to make this concrete is to reason over a small menu of predefined opponent **types** rather than the whole infinite space of strategies. Throughout the chapter we use a recurring cast, each of which serves double duty — as an opponent to play against, and as a hypothesis the detector reasons over:

Type	Behavior
Calling station	Never bets, never folds — checks when it can, calls any bet.
Rock (tight-passive)	Commits chips only with the best hand; folds everything else. The most exploitable type.
Maniac (loose-aggressive)	Always bets or calls, regardless of hand.
Nash	Balanced equilibrium play; mixes its actions; unexploitable.

7.2.2 The update, worked by hand

The mechanics are easiest to trust when you turn the crank once yourself. Suppose three candidate types bet a particular hand with probability 0.8, 0.5, and 0.1 respectively, and you start from a uniform prior $(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$. You observe a **bet**:

$$\text{posterior} \propto \left(\frac{1}{3} \cdot 0.8, \frac{1}{3} \cdot 0.5, \frac{1}{3} \cdot 0.1\right) = (0.267, 0.167, 0.033) \xrightarrow{\text{normalize}} (0.571, 0.357, 0.071).$$

Observe a **second** bet and multiply by (0.8, 0.5, 0.1) again:

$$(0.711, 0.278, 0.011).$$

The belief piles onto the bet-happy type, and it would do so *faster* the more the candidates disagree. That “how much they disagree” intuition returns in Section 6 as the reason some opponents are identified in four hands and others take two hundred.

7.2.3 Why Dirichlet, and why you only ever need the mean

Two facts from the literature shape every practical model. First, when you estimate an action distribution by counting, the natural prior is the **Dirichlet** distribution, because it is the *conjugate* prior of the multinomial: updating a Dirichlet with observed counts yields another Dirichlet, with the counts simply **added** to the prior’s pseudo-counts. No integral, no simulation — the posterior mean of action i is

$$\mathbb{E}[p_i] = \frac{\alpha_i + n_i}{\sum_j (\alpha_j + n_j)},$$

where α_i are prior pseudo-counts (which double as *smoothing*, keeping no action at exactly zero) and n_i are observed counts. This closed form *is* the “continuous” model of Section 4.

Second — and this is the single most reused result in the area — to choose an optimal response you never need the whole posterior distribution over strategies; you only need its **mean**. Your expected payoff against a *distribution* of opponent strategies equals your payoff against the single averaged strategy. That collapses an intractable integral over strategy space into “best-respond to one averaged opponent,” and it is what makes Bayesian exploitation practical. (Section 5 revisits this: responding to the mean is payoff-optimal but, as it turns out, can fail to *converge* to the opponent’s true strategy — the subtlety the 2025 work resolves.)

Read more: Ganzfried, S. & Sun, Q. (2016/2018). “Bayesian Opponent Exploitation in Imperfect-Information Games.” *IEEE CIG*. (Theorem 2.1: respond to the posterior mean.)

7.3 3. Seeing Through a Keyhole — the partial-observability problem

There is a reason opponent modeling in poker is harder than counting frequencies in rock-paper-scissors, and it deserves its own section because it shapes every model that follows.

In poker you observe **actions** — bet, call, fold — but **not** the private hand that produced them. The same “bet” can come from a monster or a bluff. The one exception is a **showdown**, where cards are revealed at the end of a hand that goes to the end without a fold. So every hand falls into one of two evidential categories:

- **Showdown** — **cause and effect**. You see both the opponent’s actions *and* the private hand behind them. You can attribute the behavior to the exact situation that produced it. This is the gold-standard observation.
- **Fold** — **effect only**. The hand ends without revealing cards. You saw *what* they did but not *what they held*, so you cannot be sure which situation to credit.

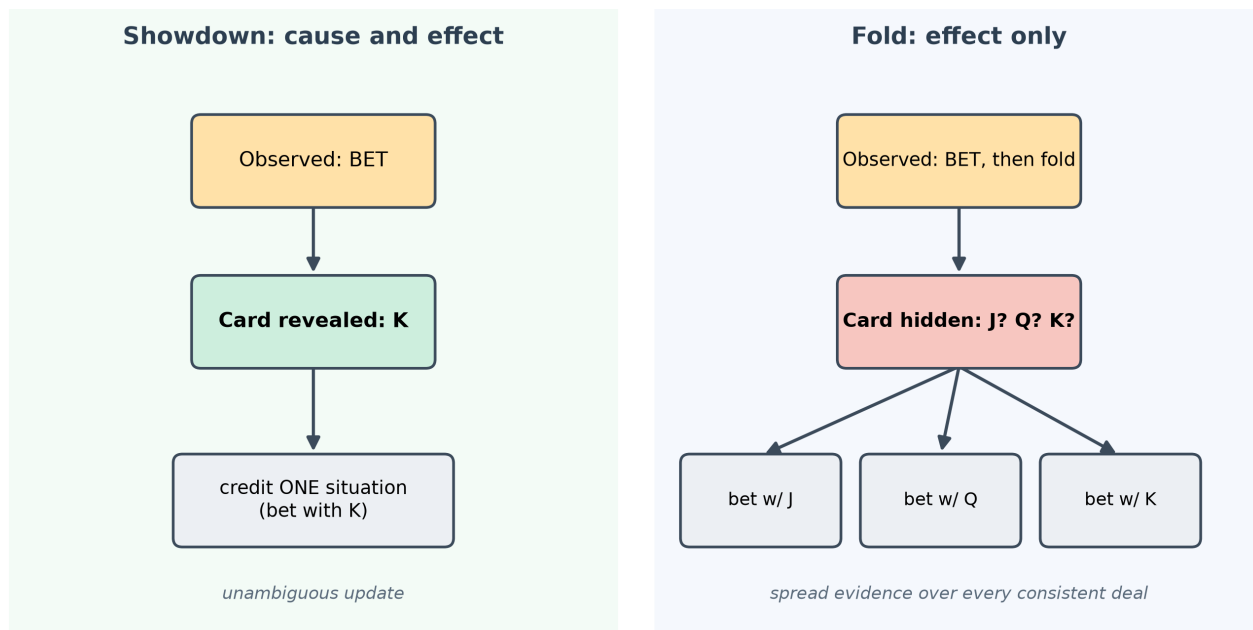


Figure 3: Partial observability: a showdown reveals the opponent’s private card, pinning the action to one situation; a fold hides it, forcing the model to spread the evidence across every hand the opponent might have held.

The consequence is a genuine theoretical limit, not an implementation nuisance: **without ever observing the opponent’s private information, you cannot learn beyond your prior**. If folds are all you ever see, more hands do not help — the evidence is fundamentally ambiguous. Showdowns are precisely the information that breaks the ambiguity.

How does a model cope with a fold? By reasoning over **all the hands the opponent might have held**. Concretely, it enumerates the deals consistent with what was actually seen, replays the opponent’s decisions under each hypothetical private card, and spreads the evidence across them — weighting each hypothesis by how plausible it is. A showdown collapses that set to a single possibility; a fold leaves it spread. This “marginalize over the hidden card” move is the technical heart of every model in the next two sections, and — not coincidentally — the formal object the consistency theory of Section 5 is built around.

7.4 4. Three Models, One Interface

There is no single “opponent model.” There is a family, trading off convergence speed, robustness to surprises, interpretability, and compute. This project implements and compares three points on that spectrum, deliberately built behind **one shared interface** — each consumes a stream of observed hands and emits a predicted opponent strategy — so that the *same* downstream best response can be applied to all three and the comparison is apples-to-apples.




Figure 4: Three models on one interface: a discrete belief over types, a free-form per-situation count, and a globally-consistent sequence-form estimate — all emit a predicted opponent strategy consumed by the same best-response step.

- **Type-based.** Maintain a belief (a probability distribution) over a fixed menu of opponent types, update it by Bayes’ rule (Section 2), and report the posterior-weighted average strategy. It converges *extremely* fast **when the opponent is in the menu**, and its output is interpretable (“80% rock, 20% maniac”). Its weakness is equally clear: if the real opponent resembles none of the types, the model has no honest way to say so (Section 6).
- **Continuous.** Drop the menu. Estimate the opponent’s action probabilities *directly at each situation* by counting what they did there, smoothed by a Dirichlet prior (Section 2). It can represent **any** strategy, including blends and novel styles, so it is the robust choice against unknown opponents. The cost is data: it learns each situation independently, with no structural sharing, so it needs many more observations — and, as Section 7 shows, that hunger has teeth on the larger game.
- **Consistent.** Estimate a single *globally consistent* strategy — one guaranteed to be a valid strategy over the whole game tree — rather than each situation in isolation. This is the most principled model and the most recent; it is important enough, and central enough to the thesis, to get its own section (Section 5).

The following table is the mental map to carry forward:

Model	Represents	Converges	Robust to out-of-menu?	Interpretable	Compute
Type-based	a few known types	fastest	no	high (named types)	cheap
Continuous	any per-situation strategy	slower (data-hungry)	yes	medium (per-situation)	cheap
Consistent	one valid global strategy	provably to the truth	yes	medium	expensive (optimization)

The obvious question — *which one should the thesis use?* — does not have a single answer, and that is itself a finding. The natural target is a **hybrid**: lean on structural priors (types) when data is sparse, and converge toward the consistent estimate as evidence accumulates. This chapter builds and measures the endpoints so that the hybrid has firm ground to stand on.

Read more: Bard, N. (2013). “Online Implicit Agent Modelling.” *AAMAS* — the explicit-vs-implicit axis that frames this taxonomy.

7.5 5. The Consistency Problem and the Sequence-Form Fix

This section is the theoretical frontier of the step and the piece the thesis most directly extends. It answers a question the earlier models quietly beg: *if a model fits the observed behavior well, have we actually recovered the opponent’s true strategy?* The surprising answer is **not necessarily** — **even with infinite data** — and there is a principled fix.

7.5.1 The flaw: fitting the mean is not the same as finding the truth

The classical Bayesian recipe responds to the posterior **mean** strategy (Section 2). Call an opponent-modeling method **consistent** if its estimate converges to the opponent’s true strategy σ^* as observations accumulate against a fixed opponent. One would hope the classical recipe is consistent. It is not.

The cleanest counterexample is rock-paper-scissors. Suppose the modeler reasons over a set of *sampled* candidate strategies and always responds to some average of them. If the true strategy $\sigma^* = (0.8, 0.1, 0.1)$ lies **outside** the region those samples can combine to, the model — being always a weighted average of the samples — simply **cannot reach it**, no matter how much data arrives. That much is intuitive. The deeper result is that the method can fail **even when the truth lies inside** the achievable region: with the true strategy sitting at the center of three samples that average to it, the posterior weight on the single best-fitting sample grows without bound relative to the others, so asymptotically the belief

collapses onto one sample rather than settling on the true mixture. Fitting the data ever more tightly, the model converges to the *wrong* strategy.




Figure 5: Why fitting the mean can miss the truth: the modeler can only ever produce points in the shaded hull of its samples. A true strategy outside the hull is unreachable; even one inside is abandoned as the posterior collapses onto a single vertex.

This matters because it undercuts the intuitive safety net “just collect more data.” Consistency is not automatic; it has to be engineered.

7.5.2 The fix: a convex program in sequence form

The remedy reformulates the estimate so that “a valid strategy” and “fits the observations” become constraints in a single, well-behaved optimization. The key change of variables is the **sequence form**: instead of per-situation action probabilities, represent the opponent’s strategy by **realization weights** y_r over action *sequences*. The legal strategies are then exactly those satisfying a set of **linear** constraints $Fy = f$, $y \geq 0$ — polynomial in the size of the game tree, rather than exponential.

Partial observability (Section 3) enters through an **observability function**: for each observed hand, the set of trajectories consistent with what you actually saw (a fold hides the opponent’s card and admits many trajectories; a showdown reveals it and admits one). The likelihood of an observation is a sum of realization weights over that consistent set. Putting a Dirichlet prior on the weights and maximizing the log-posterior gives the program

$$\max_y \sum_r (\alpha_r - 1) \log y_r + \sum_t \log \left(\sum_{r \in o(\ell_t)} q_r y_r \right) \quad \text{s.t.} \quad Fy = f, y \geq 0,$$

where $o(\ell_t)$ is the consistent-trajectory set for observation t and q_r are (normalized) chance probabilities. The decisive property: when $\alpha_r \geq 1$ this objective is **concave** — the first term is concave, the second

is a log of an affine function, and the constraints are affine — so it is a **convex** problem with **no local optima**. Any solver that finds a local optimum has found the global one. A standard **projected gradient** scheme suffices: take a gradient step, then project back onto the constraint set. In pseudocode the loop is unremarkable, which is the point —

```
initialize y feasible (F y = f, y >= 0)
repeat until converged:
  g <- gradient of the (negative) log-posterior at y
  z <- y - eta * g                                # gradient step
  y <- project z onto { y : F y = f, y >= 0 }    # convex projection
```

This method returns the **mode** of the posterior (its single most probable point), not the mean. That is the crucial trade: the mode is *consistent* — under mild conditions (the truth has positive prior density, distinct strategies produce distinguishable observations, and every opponent situation is visited infinitely often) the estimate provably converges to σ^* — whereas the payoff-optimal mean is generally intractable to compute exactly. So the design axis underneath the whole step is **mean versus mode**: payoff-optimal-but-intractable versus tractable-and-consistent.

7.5.3 What we found, honestly

We implemented this consistent estimator and verified it on Kuhn strategy recovery, where it performs as advertised: its recovered strategy sits very close to the truth (total-variation distance roughly **0.004 to 0.021**), matching or beating the continuous model’s recovery on the same game. That confirms the machinery and the theory on the small testbed.

We did **not** run it inside the online exploitation loop, nor on the larger game. The reason is the reason the literature itself flags: the estimate is the solution of an optimization that is **re-solved as observations accumulate**, and that cost grows with history — from a fraction of a second early on to many seconds per refit once tens of thousands of hands are in hand. For a real-time, hand-by-hand match that is impractical in its naive form. Rather than a gap, we treat this as the *empirical answer* to a question the step poses explicitly — *is a per-update convex solve fast enough for real-time play?* — namely **not without incremental methods** (warm-starting each solve from the last, caching the per-hand terms), which is exactly the kind of approximation Step 8 takes up. The understanding and the recovery result are what the thesis needs from this model now; the online engineering is future work.

Read more: Ganzfried, S. (2025). “Consistent Opponent Modeling in Imperfect-Information Games.” *arXiv:2508.17671*.

7.6 6. When a Model Is Confident and Wrong

Section 5 was about a subtle failure with infinite data. This section is about a blunt failure with finite data, and it is the one that most directly motivates *safe* exploitation. It is a different failure from inconsistency:

here the problem is a **misspecified menu** — an opponent the type-based model simply cannot represent — and how the model behaves when cornered.

We built a hidden opponent that is a 50/50 per-action blend of the rock and the maniac — deliberately **none** of the four candidate types. A reasonable guess is that the posterior would split its belief between the two nearest types. It does not. A product-of-likelihoods posterior concentrates on the *single best explanation*, so the belief lurches from one type to another and finally commits, hard, to **Nash** — the one candidate that assigns real probability to *both* betting and checking a middling hand, and so is never fatally contradicted. The model ends up **confident** (posterior near 1.0 on a single type) and **wrong** (that type is not what it is playing).

This exposes a trap worth stating plainly: the posterior is a **relative** quantity. “Ninety-five percent Nash” means “best fit *among these four candidates*,” not “good fit in absolute terms.” Measuring the winner’s *absolute* fit gives the game away — the same Nash hypothesis that scores well against a genuine Nash opponent scores far worse against this blend, while reporting equal confidence on paper. A small stereotype menu has no honest way to say “none of the above.”

7.6.1 How long can it stay wrong? A robustness sweep

Because confidently-wrong beliefs are dangerous, we stress-tested the phenomenon over **300 random seeds of 500 hands each**, separating two things a naive metric conflates: *slow convergence* from *falling after convergence*.

Measurement (300 seeds x 500 hands)	Result
Correct long-run winner by hand 500	300 / 300 (100%)
Ever falls back after taking the lead for good	0 / 300 (never)
A wrong type still leading past hand 100	40 / 300 (~13%)
A wrong type still leading past hand 200	14 / 300 (~5%)
Hand at which the truth locks in for good	median 23 - 90th pct 125 - worst 461

The good news: the belief eventually lands on the closest representable strategy every time, and once it locks in it never falls. The sobering news is the middle rows: in a meaningful minority of runs the model held a **confident wrong belief for well over a hundred hands** — long enough that an agent best-responding hard to it would have been adjusting to beat a phantom. This is the empirical face of the exploitation-versus-safety tension: a model can be trusted *eventually*, but “eventually” is sometimes 200 hands away, and acting on the interim belief as if it were certain is how you hurt yourself.

The honest fix is not a bigger menu — it is to stop asking “which single type?” and to lean on models that can represent blends (the continuous and consistent models of Sections 4-5), and to **scale how hard you exploit to how well-earned the read is**. That principle is the through-line into the next two sections and into Step 8.

7.7 7. From Model to Money — best response, the ceiling, and a self-inflicted leak

A model is only useful if acting on it wins. To measure that cleanly we feed each model’s estimate into the **same** exact best response and play full matches on both games, bracketing every result between two analytical yardsticks computed by hand-verified math, not simulation:

- **Nash EV** — what safe equilibrium play earns against that opponent (you cannot do worse than this if the read is useless);
- **ceiling** — the exact best-response value, the most that is extractable if you knew the opponent perfectly.

A good exploiter sits between the two and climbs toward the ceiling. First, a sanity check that the apparatus is trustworthy: the solver reproduces Kuhn’s exact equilibrium value of $-1/18$ (residual exploitability ≈ 0.002 , essentially unexploitable), and an exact best response beats uniform play by the predicted margins (+0.500 on Kuhn, +2.087 on Leduc).

Kuhn — both models reach the ceiling.

Opponent	ceiling	type-based	continuous
Rock	0.167	0.168	0.168
Maniac	0.333	0.337	0.330
Nash	-0.055	-0.055	-0.053

(A pure “always fold to a bet” opponent is even more exploitable — a ceiling of 0.975, reached to within 0.966 — but it is not one of the recurring types defined above, so it is omitted here for consistency.)

Leduc — the model class starts to matter. Leduc Hold’em adds a second betting round and a shared community card, roughly two orders of magnitude more situations than Kuhn — still exactly solvable, but large enough to separate the models.

Opponent	ceiling	type-based	continuous
Level-1	3.056	3.061	2.672
Maniac	2.177	2.199	2.038
Calling station	1.464	1.451	1.434
Rock	0.937	0.912	0.848
Nash	-0.083	-0.085	-0.175

Two results carry the message, and the second is the important one:

1. **You cannot exploit an equilibrium.** Against the Nash opponent every model earns approximately the (negative) game value and never more. The ceiling *is* essentially the game value. This confirms the exploitation elsewhere is real and not an artifact of the harness.




Figure 6: Exploitation against each Leduc opponent: the type-based model (which fits these opponents) hugs the exact best-response ceiling, while the continuous model tracks close but sits below for the hardest-to-fit types — and dips below the safe baseline against Nash.

2. **A confident-but-underfit model makes *you* exploitable.** On Leduc the continuous model *loses* to Nash — -0.175 against a -0.083 ceiling. With imperfect data over Leduc’s many situations it best-responds to a *wrong* estimate of an opponent who cannot be exploited at all, and in doing so opens a leak in its **own** play. The undersampling that merely costs it a little against exploitable opponents turns actively harmful against an unexploitable one.

This is the exploitation-versus-safety tension in a single number. When the model class fits, best response extracts the full theoretical value; when it does not, the agent both leaves money on the table *and* — more dangerously — hands some back. That second cost is precisely what Step 8’s safety mechanism (bounding the deviation from Nash by the model’s own confidence) exists to prevent.

The differences here are not seed luck. Across five seeds the standard error of per-hand profit is tiny relative to the effects: the type-based model is statistically indistinguishable from the ceiling on every type in both games, and the continuous model’s Leduc shortfall and its Nash self-leak are many standard errors wide.

Read more: Ganzfried, S. & Sandholm, T. (2015). “Safe Opponent Exploitation.” *ACM EC* — the safety half of the dial, and the anchor for Step 8.

7.8 8. Adapting to Change — non-stationary opponents

Everything so far assumed a fixed opponent. Real opponents drift and adapt, and a model that learned patiently for ten thousand hands is worse than useless the moment its subject changes style — it is now *confidently* describing a person who no longer exists. This is the open frontier the thesis is positioned to attack, and the step includes a first controlled probe.

The full adaptive agent runs the loop of Section 2 end to end: **observe** hands, **update** the model, periodically rebuild the hero strategy as a best response to the current estimate (optionally blended toward Nash for safety), and **act**. To handle change, it adds a detector: a lightweight statistical monitor on the opponent’s aggression that watches for a regime shift and, on firing, **resets** the model and drops back to safe play while it re-learns.




Figure 7: The adaptive exploitation loop: observe, update the model, best-response (blended toward Nash for safety), and act; a change-point detector can reset the model when the opponent’s style shifts.

We switched the opponent’s style at the midpoint of a 20,000-hand match and compared a **static** model (never forgets) against one with **change-point forgetting**. The result is **scenario-dependent**, and that is exactly the finding:

Game	Style switch (at the midpoint)	static (after switch)	change-point (after switch)
Kuhn	rock -> maniac	-0.116	+0.226
Leduc	rock -> maniac	+1.940	+0.525

- **On Kuhn, forgetting wins.** The strategy learned against a rock is bluff-heavy; unleashed on a maniac who calls everything, it *actively loses* (the static model goes negative). Detecting the switch and re-learning recovers to a healthy profit.

- **On Leduc, forgetting loses.** The maniac there leaks over two chips a hand, so a continuously-adapting model exploits it handsomely without any reset — while the detector, too eager, fires dozens of **false alarms** during the stable stretches (around sixty resets against a single true change), each one throwing away hard-won data and dropping to safe play.

The lesson is not “change detection is good” or “bad” — it is that **the reaction to a detected change matters as much as the detection**. A trigger-happy detector paired with a full reset-to-safe can cost more than staleness when the new opponent is exploitable enough that staleness is cheap. Cheaper, gentler responses — partial forgetting instead of a hard reset, a less nervous detector — are the clear next step, and they connect directly to Step 8’s confidence-scaled exploitation. (This experiment was run at a single seed, so read the *direction* as robust and the exact magnitudes as illustrative.)

7.9 9. Connections and Forward Pointers

What this step establishes. A good opponent model is *necessary but not sufficient* for profitable, safe adaptation. When the model class fits the opponent, best response reaches the exact extractable ceiling — the full theoretical value of modeling is realized. But under partial observability and limited data, an underfit model both leaves money on the table and, against an opponent who cannot be exploited, best-responds to a phantom and *loses*. Non-stationarity adds a second edge: forgetting a stale model helps only when staleness is actively harmful, and a careless detector’s false alarms can cost more than they save. One principle recurs across all of it: **exploitation must be scaled to how well-earned the read is**.

Backward connections. This step is the mirror image of the equilibrium work that preceded it. Where an equilibrium strategy *ignores* the opponent’s identity by design, an opponent model is the machinery that *uses* it — the two are the endpoints of the safety-exploitation spectrum. The sampling machinery that earlier steps used to traverse a game tree according to the *current* strategy reappears here to reason over an opponent’s actions according to the *inferred* model. And the idea of grouping many concrete situations into a few representative buckets — abstraction — is exactly what the type-based model does in strategy space.

Forward to Step 8 and the thesis. This step built the **sensor**; Step 8 builds the **actuator** — the mechanism that turns a model into *safe* exploitation, bounding how far you deviate from equilibrium by how much your read has earned. The continuous model’s self-inflicted leak against Nash (Section 7) is the empirical case for that mechanism; the confident-but-wrong sweep (Section 6) sets its budget; the consistency theory (Section 5) is the principled backbone the framework extends.

Open questions carried forward.

- *Consistency at real-time speed.* The consistent model is principled and accurate but its per-update convex solve does not yet fit an online match. Incremental solving (warm starts, cached terms) or Step 8’s subgame approximations are the routes to a usable real-time version.

- *Non-stationarity, properly.* All three models assume a stationary opponent for their guarantees. Confidence-scaled forgetting and better change signals are the immediate next work; a full theoretical treatment is a thesis-scale target.
- *Out-of-menu opponents.* Well-specified detection is the easy case. The interesting regime is opponents *outside* any menu — mixtures, drift, adversaries — where the non-parametric and sequence-form models should separate from the type-based one, and where an explicit “none of my types fit” signal becomes its own research question.
- *From one opponent to many.* Against several opponents at once, the jointly-optimal response is not the combination of individual best responses, and the convex guarantee of Section 5 no longer holds. Multi-opponent modeling is the bridge to the later multi-agent steps.

Read more: Shoham, Y. & Leyton-Brown, K. (2008). *Multiagent Systems*, Ch. 7 “Learning and Teaching” — the learning-in-repeated-games framing under all of the above, including the tension that your actions both *exploit* and *teach* the opponent.